

# **Soft Weighted Machine Unlearning**

---

Xinbao Qiao, Ningning Ding, Yushi Cheng, Meng Zhang

<https://arxiv.org/abs/2505.18783>

Consider a prediction task (regression problem) with:

- Input space  $\mathcal{X}$ ;
- Output space  $\mathcal{Y}$ ;
- Training set  $\mathcal{D}^n = \{z_i = (\mathbf{x}_i, y_i)\}_{\{i=1\}}^n \subset \mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ ;
- Parameter  $\theta \in \Theta := \mathbb{R}^d$ ;
- $f(\theta; \mathbf{x})$  estimator of  $y|\mathbf{x}$ ;
- $\ell : \mathcal{Z} \times \Theta \rightarrow \mathbb{R}$  loss function (e.g.,  $\ell(z, \theta) \mapsto \|y - f(\theta; \mathbf{x})\|^2$ ).

We seek the empirical risk minimizer:

$$\hat{\theta} = \arg \min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(z_i, \theta).$$

For some problems (data poisoning, fairness, robustness, etc.), it is desirable to unlearn the contribution of certain training data.

### Challenges:

- Retraining the model from scratch is computationally expensive;
- The removal may decrement the performance of the model on the remaining data.

***Solution:*** Use influence functions with soft weights.

Hard



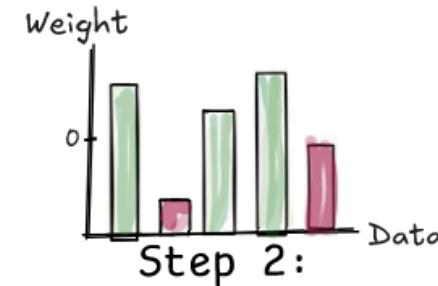
Step1:

### Influence Evaluation

1. Compute each sample's influence on the target task

Soft

1. Compute each sample's influence on the target task
2. Compute each sample's influence on the utility.



### Data Reweighting

Sample z: ...

Weight  $\epsilon$ : -1 -1 0 0 0

Select the top-k detrimental samples on target task, and set weights to -1 (unlearn).

Sample z: ...

Weight  $\epsilon$ : -1.2 -0.3 0.1 0.8 1.1

Optimize a set of weights based on the influence on the utility and the target task.

# Influence Functions

Given  $z_j \in \mathcal{D}^n$  and  $\varepsilon_j \in \mathbb{R}$ , we can define:

$$\hat{\theta}(z_j, \varepsilon_j) = \arg \min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(z_i, \theta) + \varepsilon_j \ell(z_j, \theta).$$

The *influence function* relative to  $z_j$  is:

$$\mathcal{I}(z_j) = \left. \frac{d}{d\varepsilon_j} \hat{\theta}(z_j, \varepsilon_j) \right|_{\varepsilon_j=-1}$$

**Interpretation:** It indicates how the training error would change if we removed the training data  $z_j$ .

1. If  $\ell$  is at least  $C^2$ , we can write the influence function in a closed form:

$$\mathcal{I}(z_j, -1) = -\frac{1}{n} H_{\hat{\theta}}^{-1} \nabla_{\theta} \ell(z_j, \hat{\theta})$$

where  $H$  is the Hessian matrix of the empirical risk at  $\hat{\theta}$ .

2. The influence function is a *first-order approximation* of the change in the estimator when we remove  $z_j$  from the training set.

In fact, if we do the *Taylor expansion* of  $\hat{\theta}$  at  $\varepsilon_j = -1$ , we have:

$$\hat{\theta} = \hat{\theta}(z_j, -1) + \mathcal{I}(z_j) + O(\varepsilon_j^2).$$

We will consider the influence on these 3 metrics (negative influence is better):

- **Utility:**  $\mathcal{I}_{\text{util}}(z_j, -1) = \sum_{z \in \mathcal{T}} \nabla_{\theta} \ell(z, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_{\theta} \ell(z, \hat{\theta})$ , where  $\mathcal{T}$  is the validation set;
- **Fairness:**  $\mathcal{I}_{\text{fair}}(z_j, -1) = \nabla_{\theta} f_{\text{fair}}(\mathcal{T}, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_{\theta} \ell(z, \hat{\theta})$ , with  $f_{\text{fair}}$  a fairness metric (e.g., demographic parity);
- **Robustness:**  $\mathcal{I}_{\text{robust}}(z_j, -1) = \sum_{\tilde{z} \in \tilde{\mathcal{T}}} \nabla_{\theta} \ell(\tilde{z}, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_{\theta} \ell(z, \hat{\theta})$ , for a perturbed dataset  $\tilde{\mathcal{T}}$ .

---

**Algorithm 1:** Soft-Weighted Unlearning Framework

---

**Input:** Model  $\hat{\theta}$ , Training Dataset  $\mathcal{D}$ , Validationa and Testing Dataset  $\mathcal{T}$ , Adversarial Samples  $\tilde{z} \in \tilde{\mathcal{T}}$

**# Step 1: Influence Evaluation.**

**for** each sample  $z_i \in \mathcal{D}$  **do**

Evaluate influence of  $z_i$  on validation set;

Utility:  $\mathcal{I}_{\text{util}}(z_i; -1) \leftarrow$  Equation (3).

Fairness:  $\mathcal{I}_{\text{fair}}(z_i; -1) \leftarrow$  Equation (4).

Robustness:  $\mathcal{I}_{\text{robust}}(z_i; -1) \leftarrow$  Equation (5).

**end**

**# Step 2: Weights Optimization.**

Weights  $\{\epsilon_i^*\}_{i=1}^n \leftarrow$  Equation (7)

**# Step 3: Model Correction.**

**if**  $f \leftarrow f_{\text{fair}}(\mathcal{T}; \theta)$  or  $\sum_{\tilde{z} \in \tilde{\mathcal{T}}} \nabla_{\theta} \ell(\tilde{z}; \theta) \leq \delta$  **then**

$\theta \leftarrow$  Equation (8) or Other Unlearning Algorithms

**end**

**Output:**  $\theta$

---

Instead of a binary decision (keep or remove), we want to assign a customized weight to each training data.

Namely, we want to find  $\varepsilon^* = (\varepsilon_1, \dots, \varepsilon_n)$  and use it for the one-shot updating rule:

$$\hat{\theta}(\mathcal{D}, \varepsilon^*) \approx \hat{\theta} + \varepsilon^* H_\theta^{-1} \left( \nabla_{\theta} \ell(z_1, \hat{\theta}), \dots \nabla_{\theta} \ell(z_n, \hat{\theta}) \right)^\top.$$

I.e., we update each training data with a re-scaled loss function

$$\varepsilon_i \mathcal{J}(z_i, -1) = \mathcal{J}(z_i, \varepsilon_i).$$

We need weights to be easy to retrieve, since they must be computed for each training data.

They are defined as the solution of the following optimization problem:

$$\min_{\boldsymbol{\varepsilon} \in \mathbb{R}^n} \sum_{i=1}^n \mathcal{J}_{\text{fair}}(z_i, \varepsilon_i) + \lambda \|\boldsymbol{\varepsilon}\|_2^2$$

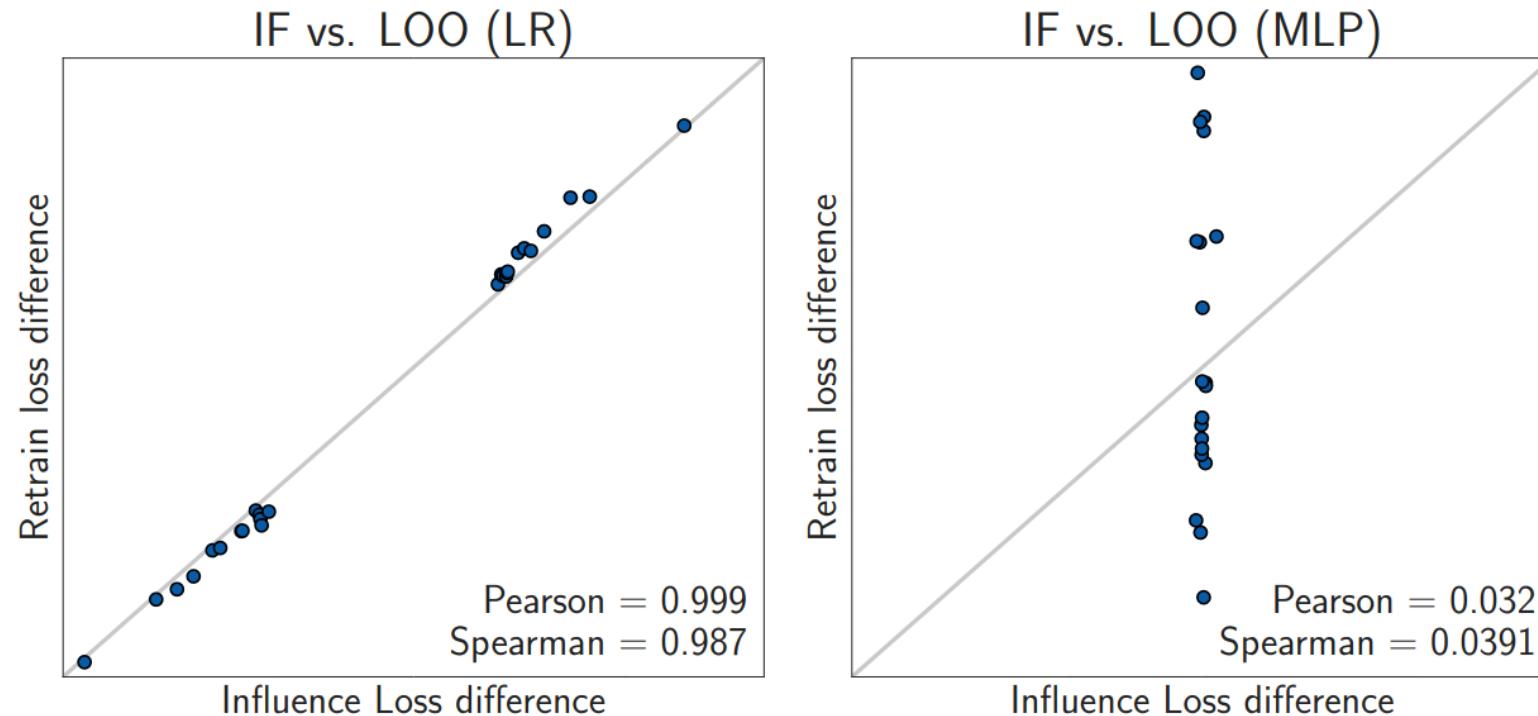
(quadratic problem) subject to:

$$\sum_{i=1}^n \mathcal{J}_{\text{fair}}(z_i, \varepsilon_i) \geq -f_{\text{fair}}(\mathcal{T}; \hat{\theta}), \quad \sum_{i=1}^n \mathcal{J}_{\text{util}}(z_i, \varepsilon_i) \leq 0$$

- We assume to know  $\hat{\theta}$  so that  $H_{\hat{\theta}}$  is PD, thus invertible. In practice we only have an approximation;
- Computing the inverse of the hessian is expensive. We must use approximations;
- Such approximations don't work well for neural networks.

# Bad approximation

11 / 15



[“If Influence Functions are the Answer, Then What is the Question?”, J. Bae, N. Ng, A. Lo, M. Ghassemi, R. Grosse, 2022]

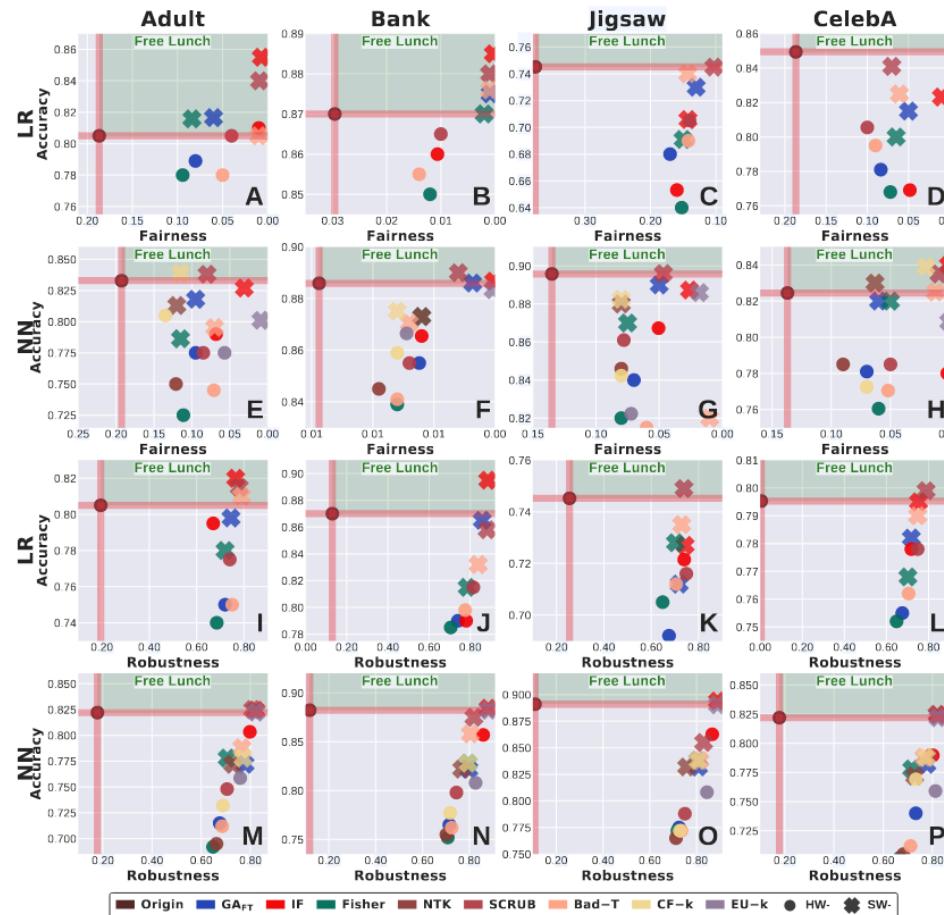
1. Add a regularization term to the loss function. This way,  $H_{\hat{\theta}}$  is guaranteed to be PD.
2. Approximate the hessian with  $\sigma I$ . This way, we have:

$$\hat{\theta}(z_j, \varepsilon_j^*) - \hat{\theta} \approx \varepsilon_j^* \frac{\sigma}{n} \nabla_{\theta} \ell(z_j, \hat{\theta}).$$

If we see  $\frac{\sigma}{n}$  as a learning rate, we can index it by time and use also others unlearning algorithms instead of IFs to do the update in an iterative fashion(?).

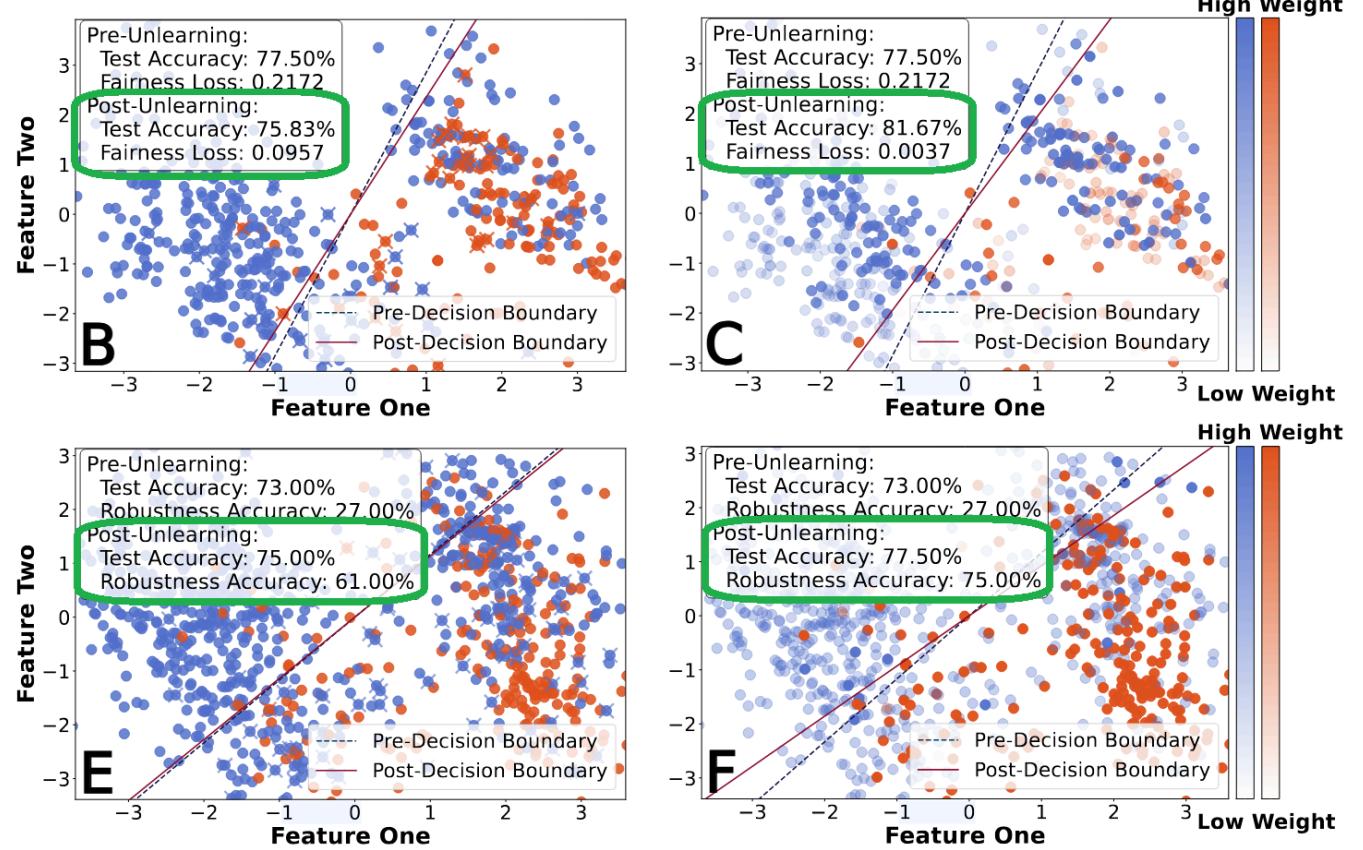
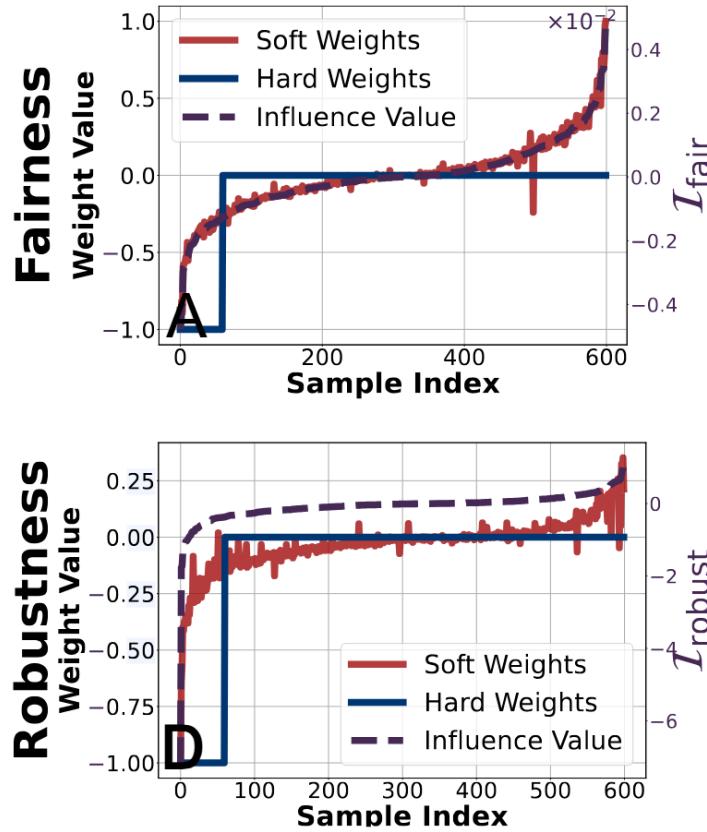
# Numerical evidence

13 / 15



# Numerical evidence

14 / 15



## Take-away

Using some smart approximations, we can adapt all the existing unlearning algorithms to the soft weighted setting.

This family of methods performs better than the HW version for non-privacy applications and can also improve the performance of the model on the remaining data.

[“Soft Weighted Machine Unlearning”, X. Qiao, N. Ding, Y. Cheng, M. Zhang, 2025]